



Movie Recommendation System

A Machine Learning Journey

SLIDE 1: Title Slide



Building a Smart Movie Recommender

Using Machine Learning to Suggest Movies You'll Love

Team: Daniel

Date: December 2024

"Imagine Netflix knowing exactly what you want to watch tonight — that's what we built!"

SLIDE 2: The Problem We're Solving



The Problem: Too Many Choices!

Have you ever...

- Spent **30 minutes** scrolling through movies and gave up?
- Watched something random and **hated it**?
- Wished someone could just **pick for you**?

The Numbers:

Platform	Movies Available
Netflix	15,000+
Amazon Prime	24,000+
All Streaming	100,000+

🎯 Our Goal:

| Build a system that learns YOUR taste and suggests movies YOU'LL actually enjoy.

Think of it like: Having a friend who knows exactly what movies you like!

SLIDE 3: What is Machine Learning?

🤖 Machine Learning Explained Simply

Imagine Teaching a Child:

Traditional Programming:

You → Tell computer EXACT rules → Computer follows rules

Machine Learning:

You → Show computer EXAMPLES → Computer LEARNS patterns

Real-Life Example:

How Humans Learn	How Our Model Learns
"Mom showed me 100 cat photos"	"We showed it 25 million movie ratings"
"Now I can recognize any cat"	"Now it can predict what you'll like"

Key Point:

We don't tell the computer "recommend action movies to young men."
Instead, we show it millions of examples, and it **discovers patterns on its own!**

SLIDE 4: Our Data — Where It All Starts

The MovieLens Dataset

What is MovieLens?

A research project by University of Minnesota where **real people** rated **real movies**.

Our Dataset Contains:

What	How Much
 Users	162,000 people
 Movies	62,000 films
 Ratings	25 MILLION ratings
 Time Span	1995 - 2019

Why This Dataset?

- Real data** — Not fake or generated
- Large enough** — ML needs lots of examples
- Free & trusted** — Used by researchers worldwide

Common Misunderstanding:

"More data = better model" is **mostly true**, but the data must be **quality data**.

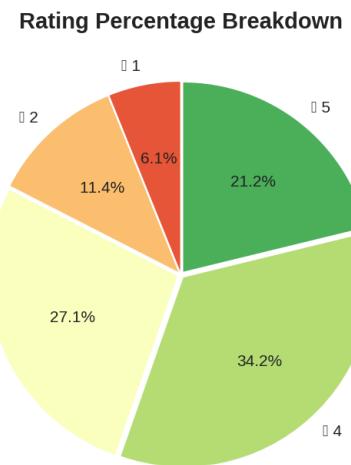
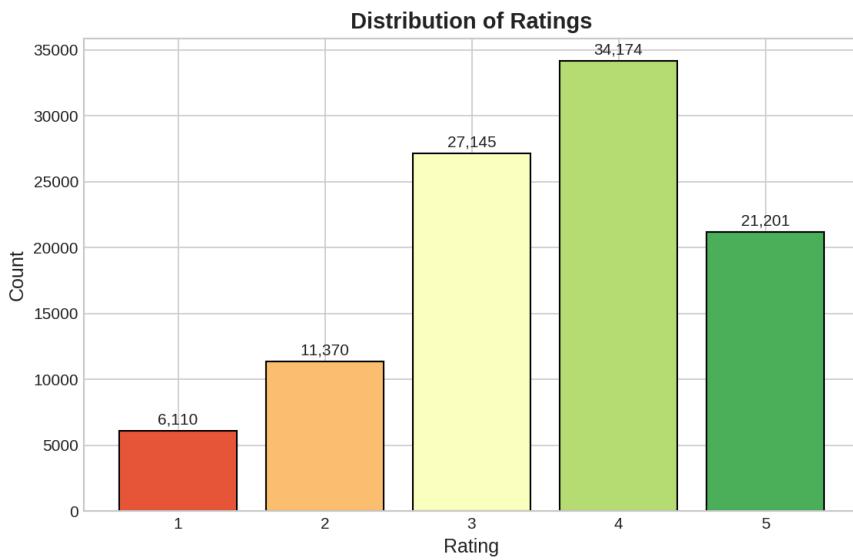
25 million ratings from real users is much better than 100 million fake ratings!

SLIDE 5: Exploring Our Data (Part 1)

🔍 What Did We Find?

Rating Distribution

"How do people typically rate movies?"



📊 PLOT: 01_rating_distribution.png

What you'll see:

- Most ratings are 3-4 stars
- People rarely give 1 star (they just don't watch bad movies!)
- 4 stars is the most common rating

💡 Insight:

People are **generous raters!** The average rating is about **3.5 stars**.

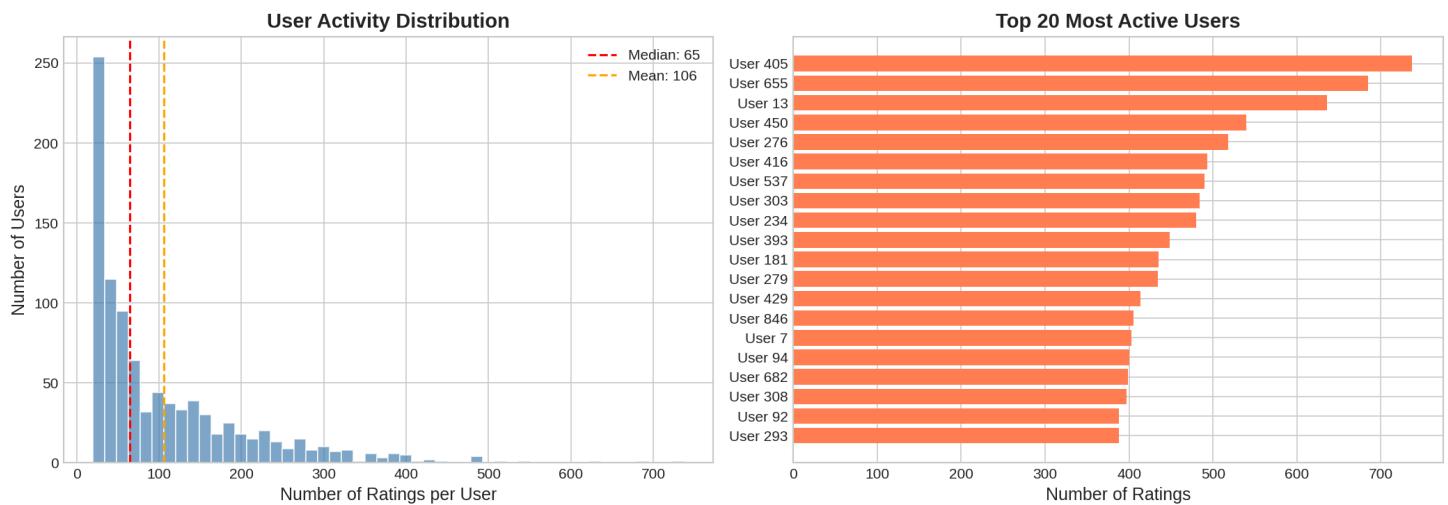
This tells us: If our model predicts 3.5 for an unknown movie, it's already making a reasonable guess!

SLIDE 6: Exploring Our Data (Part 2)



User Behavior Patterns

How Active Are Users?



▀ PLOT: 02_ratings_per_user.png

What you'll see:

- Most users rated 20-50 movies
- Some "super users" rated 1000+ movies
- Many users rated very few movies (the "cold start" problem)



Why This Matters:

- Users with **many ratings** → Easy to predict their taste
- Users with **few ratings** → Harder to predict (we need creative solutions!)



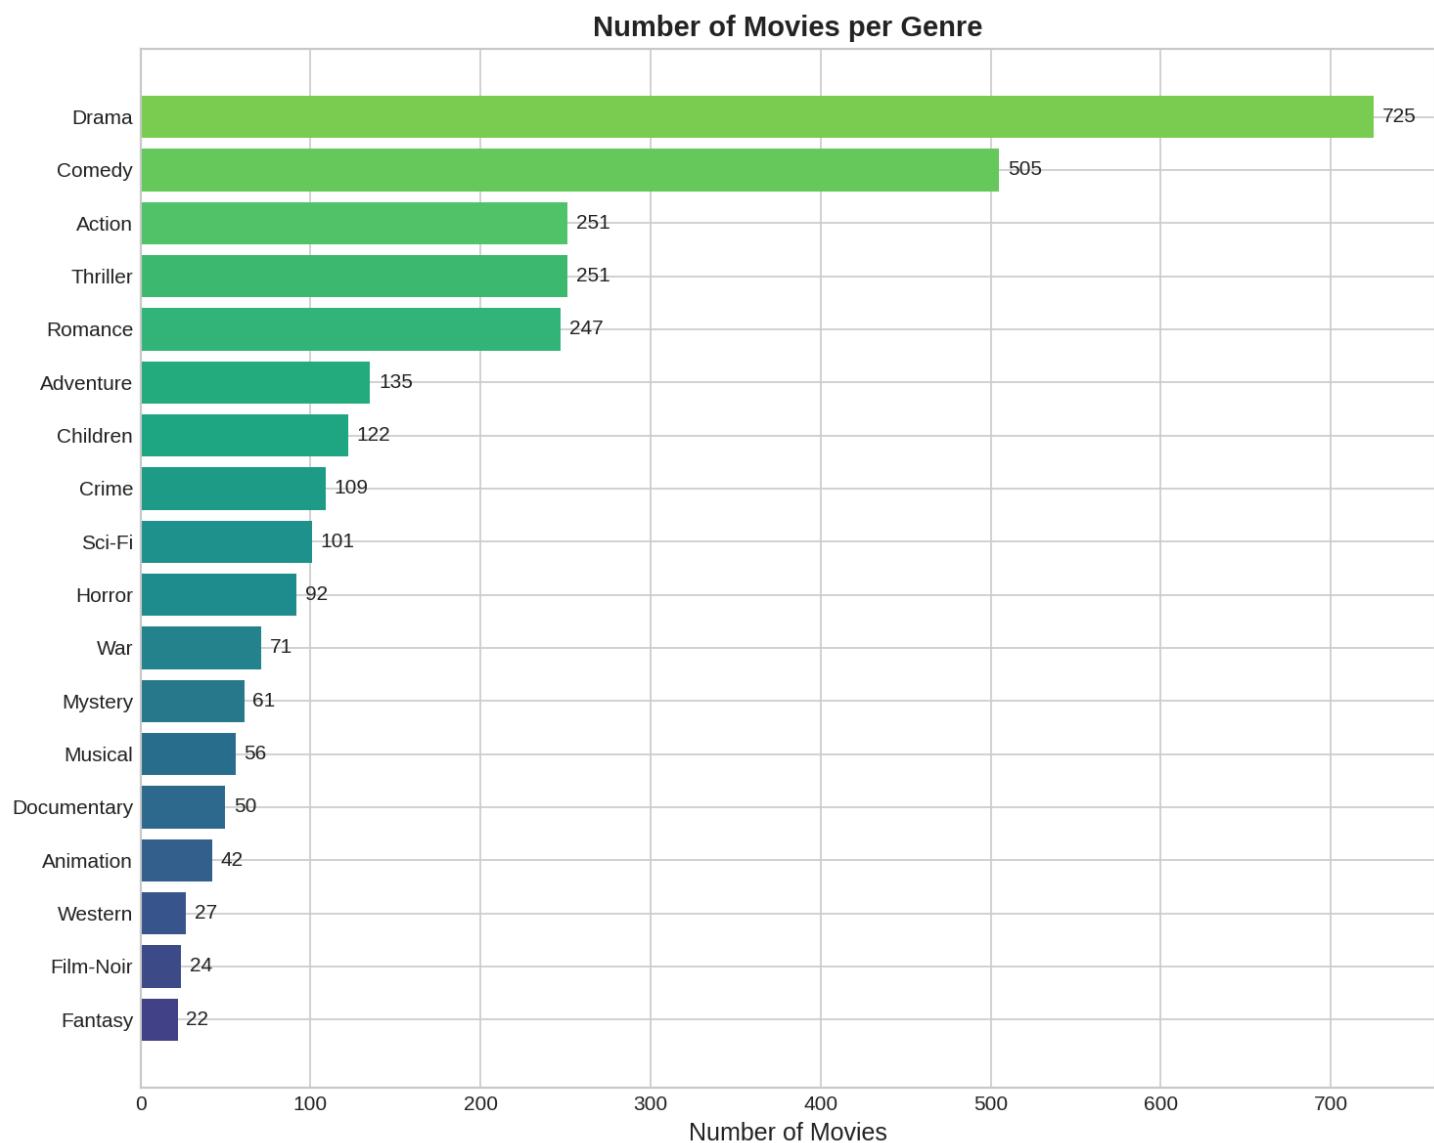
The Cold Start Problem:

When a **new user** joins with **zero ratings**, how do we recommend anything?
(We'll solve this later with our hybrid approach!)

SLIDE 7: Exploring Our Data (Part 3)

🎭 What Genres Are Popular?

Genre Distribution



📊 PLOT: 03_genre_distribution.png

What you'll see:

- Drama is the most common genre
- Comedy is second
- Some genres (Film-Noir, IMAX) are very rare

Insight:

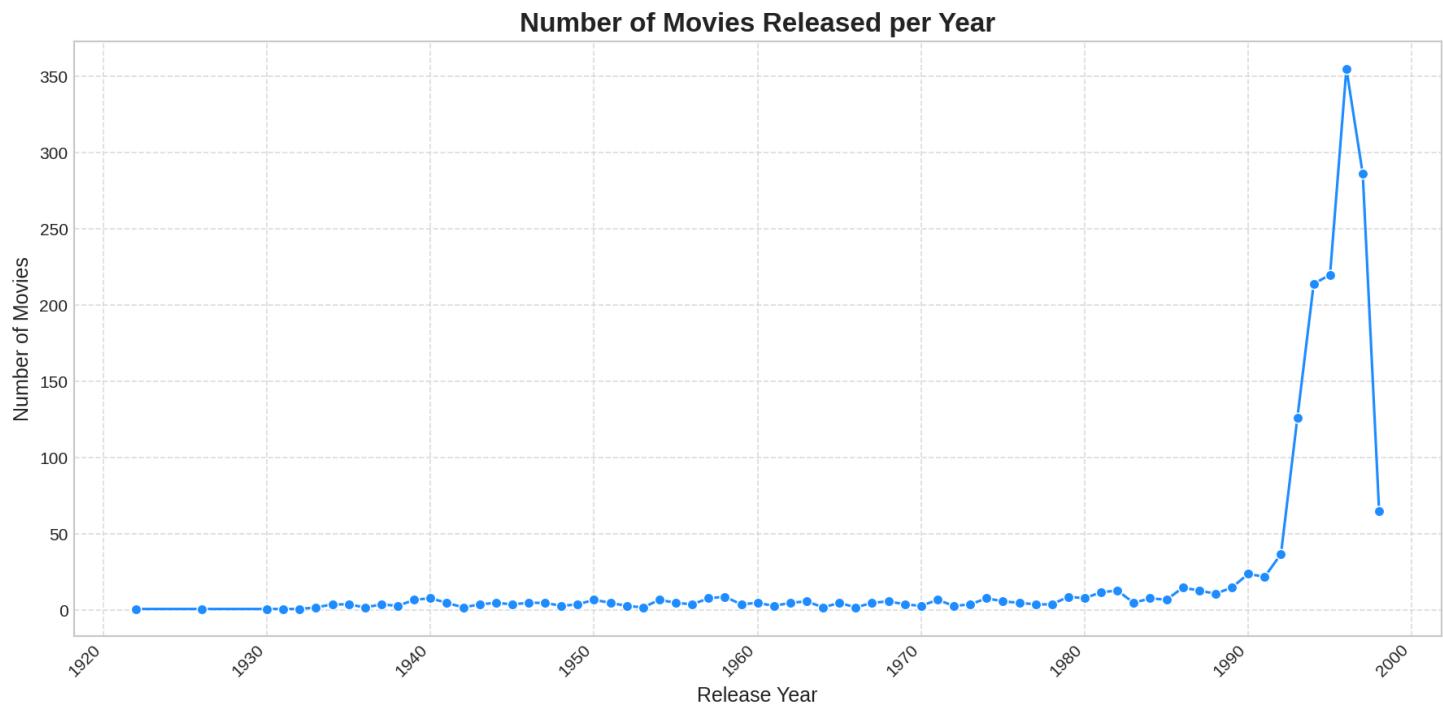
We have **unbalanced data** — lots of dramas, few documentaries.

The model might get better at recommending dramas simply because it has more examples!

SLIDE 8: Exploring Our Data (Part 4)

Movies Over Time

Release Year Distribution



 PLOT: 04_movies_by_year.png

What you'll see:

- Most movies are from 1900-2000
- Older classics (before 1970) are rare
- Recent years have LOTS of movies

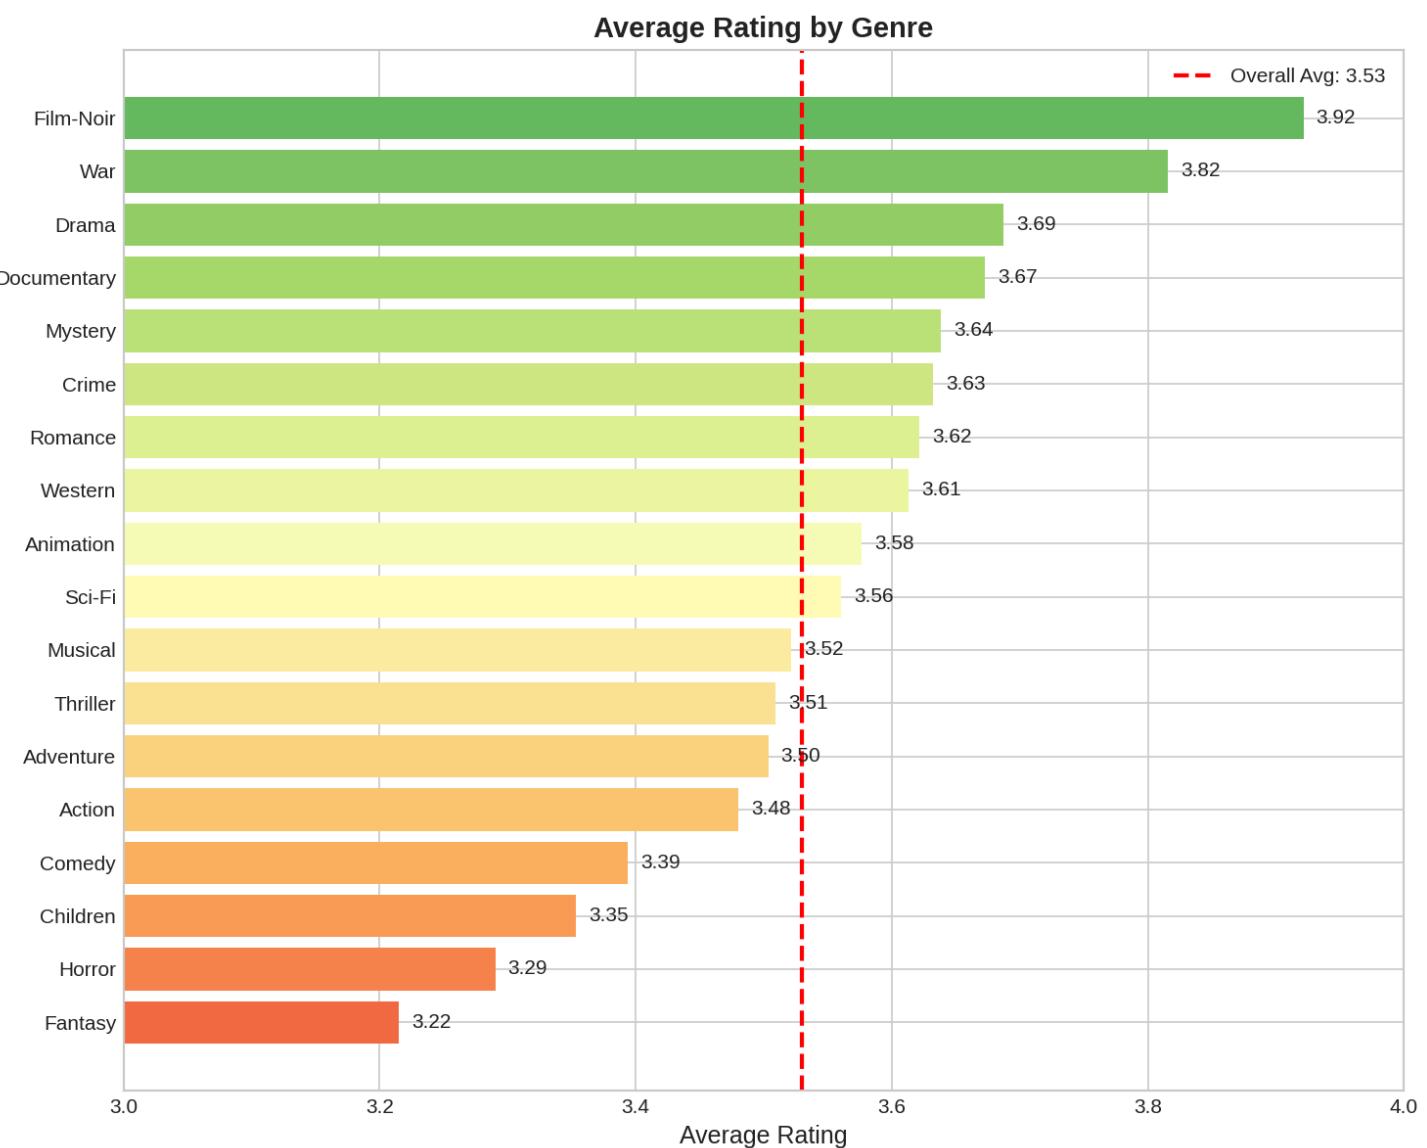
Why This Matters:

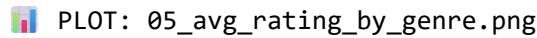
- Newer movies have **more ratings** (more people watched them)
- Classic films have **fewer but passionate** ratings
- Our model sees more recent movies, so it might be biased toward them

SLIDE 9: Exploring Our Data (Part 5)

Rating Patterns

Average Rating by Genre





What you'll see:

- Film-Noir and War movies have highest average ratings
- Horror movies have lowest average ratings
- This doesn't mean Horror is "bad" – just rated differently!

⚠ Common Misunderstanding:

"Horror has low ratings, so we shouldn't recommend horror?"

WRONG! A horror fan who rates horror movies 4 stars is a HAPPY horror fan.

The model learns **personal preferences**, not just what's "highest rated."

SLIDE 10: Preparing the Data

🧹 Data Preprocessing

Why Clean Data?

Raw Data → Like a messy room

Clean Data → Like an organized library

Machine learning models learn better from organized data!

What We Did:

Step	What	Why
1	Removed movies with < 5 ratings	Not enough info to learn from
2	Removed users with < 5 ratings	Not enough info about their taste
3	Extracted year from title	"Toy Story (1995)" → Year: 1995
4	Split genres into list	"Action Comedy" → ["Action", "Comedy"]

Step	What	Why
5	Normalized ratings	Made all ratings comparable

💡 **Analogy:**

Before teaching someone to cook, you organize ingredients.

Before teaching ML, you organize data!

SLIDE 11: Splitting the Data

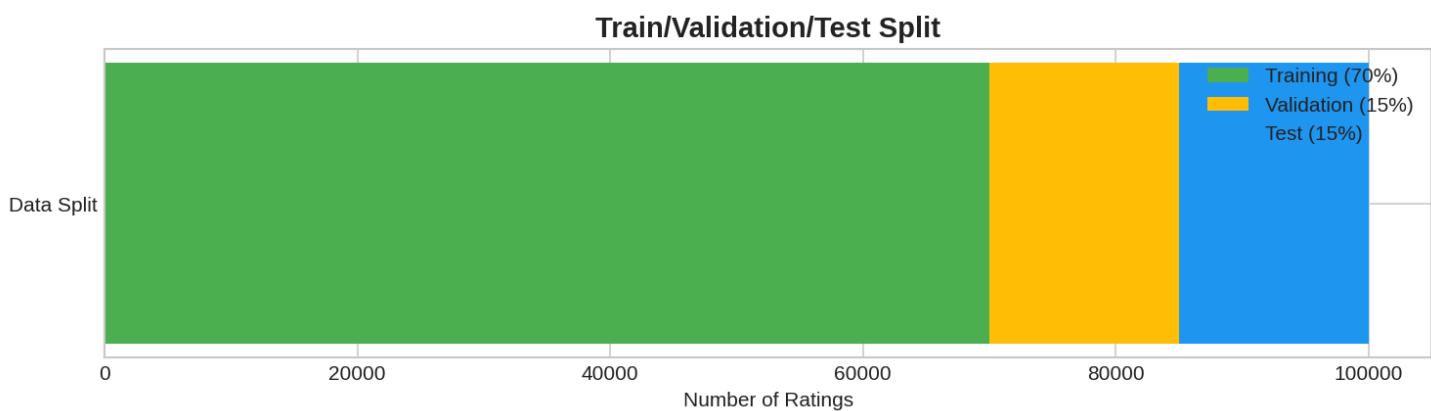
🔗 **Train, Validation, and Test Sets**

Why Split Data?

Imagine studying for an exam:

- **Study material** = Training data (learn from this)
- **Practice test** = Validation data (tune your approach)
- **Final exam** = Test data (honest evaluation)

Our Split:



 PLOT: 08_data_split.png

Training: 70% (17.5 million ratings)
Validation: 15% (3.75 million ratings)
Test: 15% (3.75 million ratings)

Critical Rule:

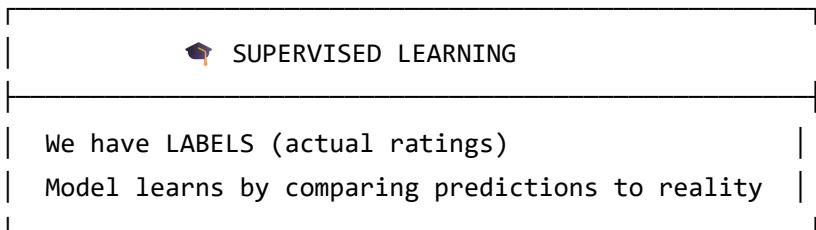
The model **NEVER sees** test data during training!
This ensures our final score is **honest and realistic**.

If we cheated and trained on test data, we'd get fake-good scores that don't reflect real performance.

SLIDE 12: Supervised Learning

Our Learning Approach

What Type of Learning?



The Learning Loop:

1. Model sees: User 42 + Movie 100
2. Model predicts: 3.5 stars
3. Actual rating: 4.0 stars
4. Error: 0.5 stars off
5. Model adjusts to be more accurate
6. Repeat millions of times!

Why "Supervised"?

Learning Type	Has Labels?	Example
Supervised <input checked="" type="checkbox"/>	Yes — actual ratings	"This user gave 4 stars"
Unsupervised	No labels	"Group similar movies"
Reinforcement	Rewards/penalties	"Game AI learns to win"

Our Case:

We have 25 million "correct answers" (actual ratings).
The model learns by trying to match these answers!

SLIDE 13: Our First Approach — Collaborative Filtering

Finding Similar Users

The Idea:

"People who agreed in the past will agree in the future."

Simple Example:

You rated: Inception ★★★★★ Titanic ★★★★
Sarah rated: Inception ★★★★★ Titanic ★★★★ Interstellar ★★★★★

You haven't seen Interstellar.

Sarah is similar to you.

Sarah loved Interstellar.

→ Recommend Interstellar to you!

Why "Collaborative"?

Many users "collaborate" unknowingly by rating movies.

Their combined ratings reveal patterns no single user could!

SLIDE 14: The SVD Algorithm



How Collaborative Filtering Works

The Challenge:

We have a HUGE table of users × movies, mostly empty:

	Movie1	Movie2	Movie3	Movie4	...	Movie62000
User1	4	?	?	2	...	?
User2	?	5	?	?	...	3
User3	3	?	4	?	...	?
...						
User162000	?	?	?	?	...	?

Problem: 99% of cells are empty (users haven't rated most movies)!

SVD Solution:

SVD = Singular Value Decomposition

Big Sparse Matrix → Two Smaller Dense Matrices
(mostly empty) (filled with learned patterns)

Analogy:

Imagine describing a person with just 50 characteristics instead of 1000.

SVD finds the **50 most important patterns** in user preferences!

SLIDE 15: SVD Visualized

Matrix Factorization

What SVD Does:

![SVD Concept Placeholder]

 PLOT: 11_svd_concept.png

Original Matrix (Users × Movies):

4	?	?	2	?	...		162,000 users
?	5	?	?	3	...		x
3	?	4	?	?	...		62,000 movies

Becomes:

User Matrix \times Movie Matrix

$\begin{bmatrix} u1: [0.2, 0.8, \dots] \\ u2: [0.9, 0.1, \dots] \\ \dots \end{bmatrix}$	\times	$\begin{bmatrix} m1: [0.3, 0.5, \dots] \\ m2: [0.7, 0.2, \dots] \\ \dots \end{bmatrix}$
$(162K \times 100)$		$(100 \times 62K)$

The Matrix Factorization Process:

ORIGINAL MATRIX (Sparse - 99% empty)

	Movie1	Movie2	Movie3	...	Movie62K
User1	4	?	?	...	?
User2	?	5	?	...	3
User3	3	?	4	...	?
...					
User162K	?	?	?	...	?

↓ SVD Decomposition ↓

USER MATRIX (162K × 100) × MOVIE MATRIX (100 × 62K)

$$\begin{bmatrix} | & u1: [0.2, 0.8, \dots] \\ | & u2: [0.9, 0.1, \dots] \\ | & u3: [0.5, 0.6, \dots] \\ | & \dots \\ | & u162K: [...] \end{bmatrix} \times \begin{bmatrix} | & m1: [0.3, 0.5, \dots] \\ | & m2: [0.7, 0.2, \dots] \\ | & m3: [0.4, 0.3, \dots] \\ | & \dots \\ | & m62K: [...] \end{bmatrix}$$

(Dense - filled)

(Dense - filled)

What These 100 Features Represent:

The 100 latent features discovered by SVD capture **hidden patterns**:

- Feature 1: Sci-fi preference
- Feature 2: Comedy enjoyment
- Feature 3: Preference for classics
- (... 97 more abstract patterns)

The model learns these patterns **automatically** — we never explicitly tell it what they mean!

Making Predictions:

To predict User1's rating for Movie5:

```
Rating = User1_Vector • Movie5_Vector  
= [0.2, 0.8, 0.5, ...] • [0.4, 0.6, 0.2, ...]  
= (0.2×0.4) + (0.8×0.6) + (0.5×0.2) + ...  
= 3.7 stars ★★★
```

Prediction:

```
Rating(User1, Movie5) = User1_vector • Movie5_vector  
= [0.2, 0.8, ...] • [0.4, 0.6, ...]  
= 3.7 stars (predicted!)
```

SLIDE 16: Training the SVD Model

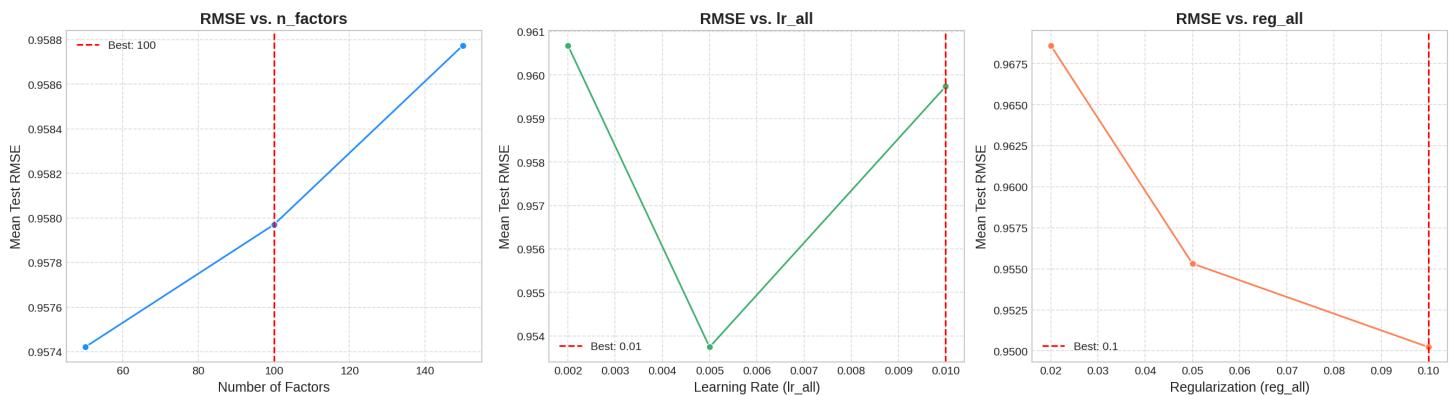
⚙️ Hyperparameter Tuning

What Are Hyperparameters?

Settings WE choose before training (the model can't learn these):

Parameter	What It Controls	Our Choice
n_factors	How many patterns to learn	100
n_epochs	How many times to see all data	20
lr_all	How fast to learn	0.005
reg_all	How much to prevent overfitting	0.02

Finding the Best Settings:





We tried many combinations:

- n_factors: [50, 100, 150, 200]
- lr_all: [0.002, 0.005, 0.01]
- reg_all: [0.01, 0.02, 0.05]

Best combination had lowest validation error!

⚠️ Overfitting Warning:

If we use too many factors, the model **memorizes** training data instead of **learning patterns**. Like memorizing exam answers instead of understanding the subject!

SLIDE 17: Our Second Approach — Content-Based Filtering

🎬 Finding Similar Movies

The Idea:

"If you liked movie A, you'll like movies similar to A."

What Makes Movies "Similar"?

Feature	Example
Genre	Action, Comedy, Drama
Director	Christopher Nolan, Spielberg
Actors	Tom Hanks, Meryl Streep
Keywords	"space", "heist", "romantic"

Feature	Example
Year	1990s, 2010s

Advantage Over Collaborative:

- Works for **new movies** with zero ratings!
- Just compare features to movies the user already liked.

SLIDE 18: TF-IDF Explained

Turning Movies into Numbers

The Problem:

Computers don't understand "Drama" or "Action" — they only understand **numbers!**

TF-IDF Solution:

TF-IDF = Term Frequency × Inverse Document Frequency

Simple Explanation:

- TF: How often a word appears in THIS movie's description
- IDF: How rare is this word across ALL movies

TF-IDF = Common in this movie + Rare overall = IMPORTANT!

Example:

Word	TF (in Inception)	IDF (all movies)	TF-IDF
"the"	Very common	Very common	LOW (useless)
"dream"	Common	Rare	HIGH (meaningful!)
"heist"	Common	Rare	HIGH (meaningful!)

Result:

Each movie becomes a **vector of numbers** representing its content:

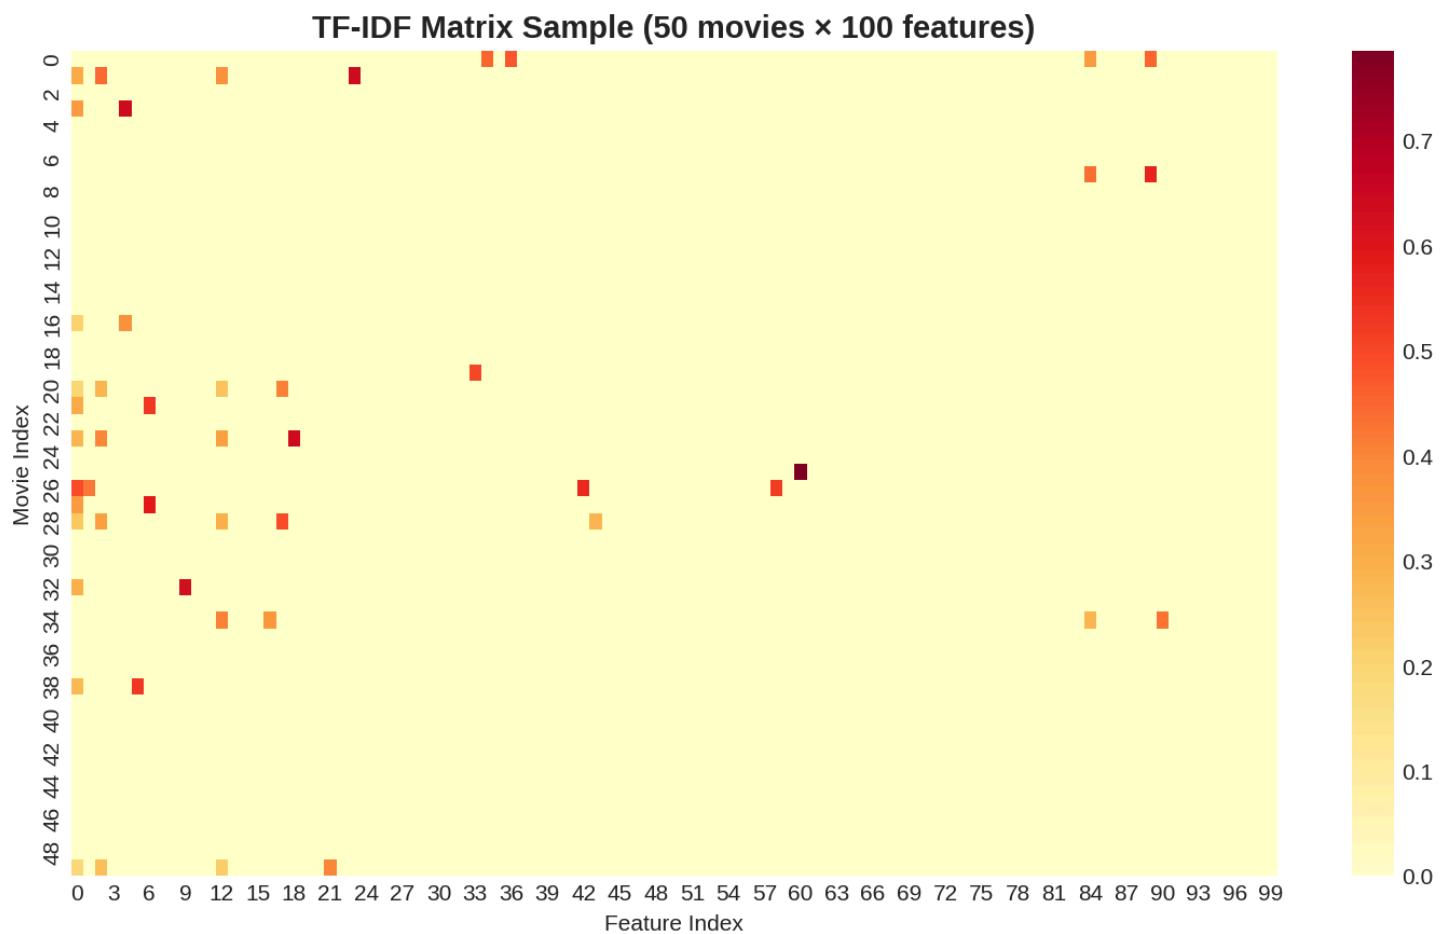
Inception = [0.0, 0.8, 0.9, 0.2, 0.7, ...] (1000 numbers)

SLIDE 19: Cosine Similarity

📐 Measuring Movie Similarity

How Similar Are Two Movies?

We compare their TF-IDF vectors using **cosine similarity**:



 PLOT: 09_tfidf_matrix.png

Cosine Similarity = $\cos(\text{angle between vectors})$

Result ranges from:

- 1.0 = Identical (same direction)
- 0.0 = Unrelated (perpendicular)
- -1.0 = Opposite (never happens with TF-IDF)

Example:

Similarity(Inception, Interstellar) = 0.85 (very similar!)

Similarity(Inception, Frozen) = 0.12 (not similar)

Why Cosine?

It measures **angle**, not **length**.

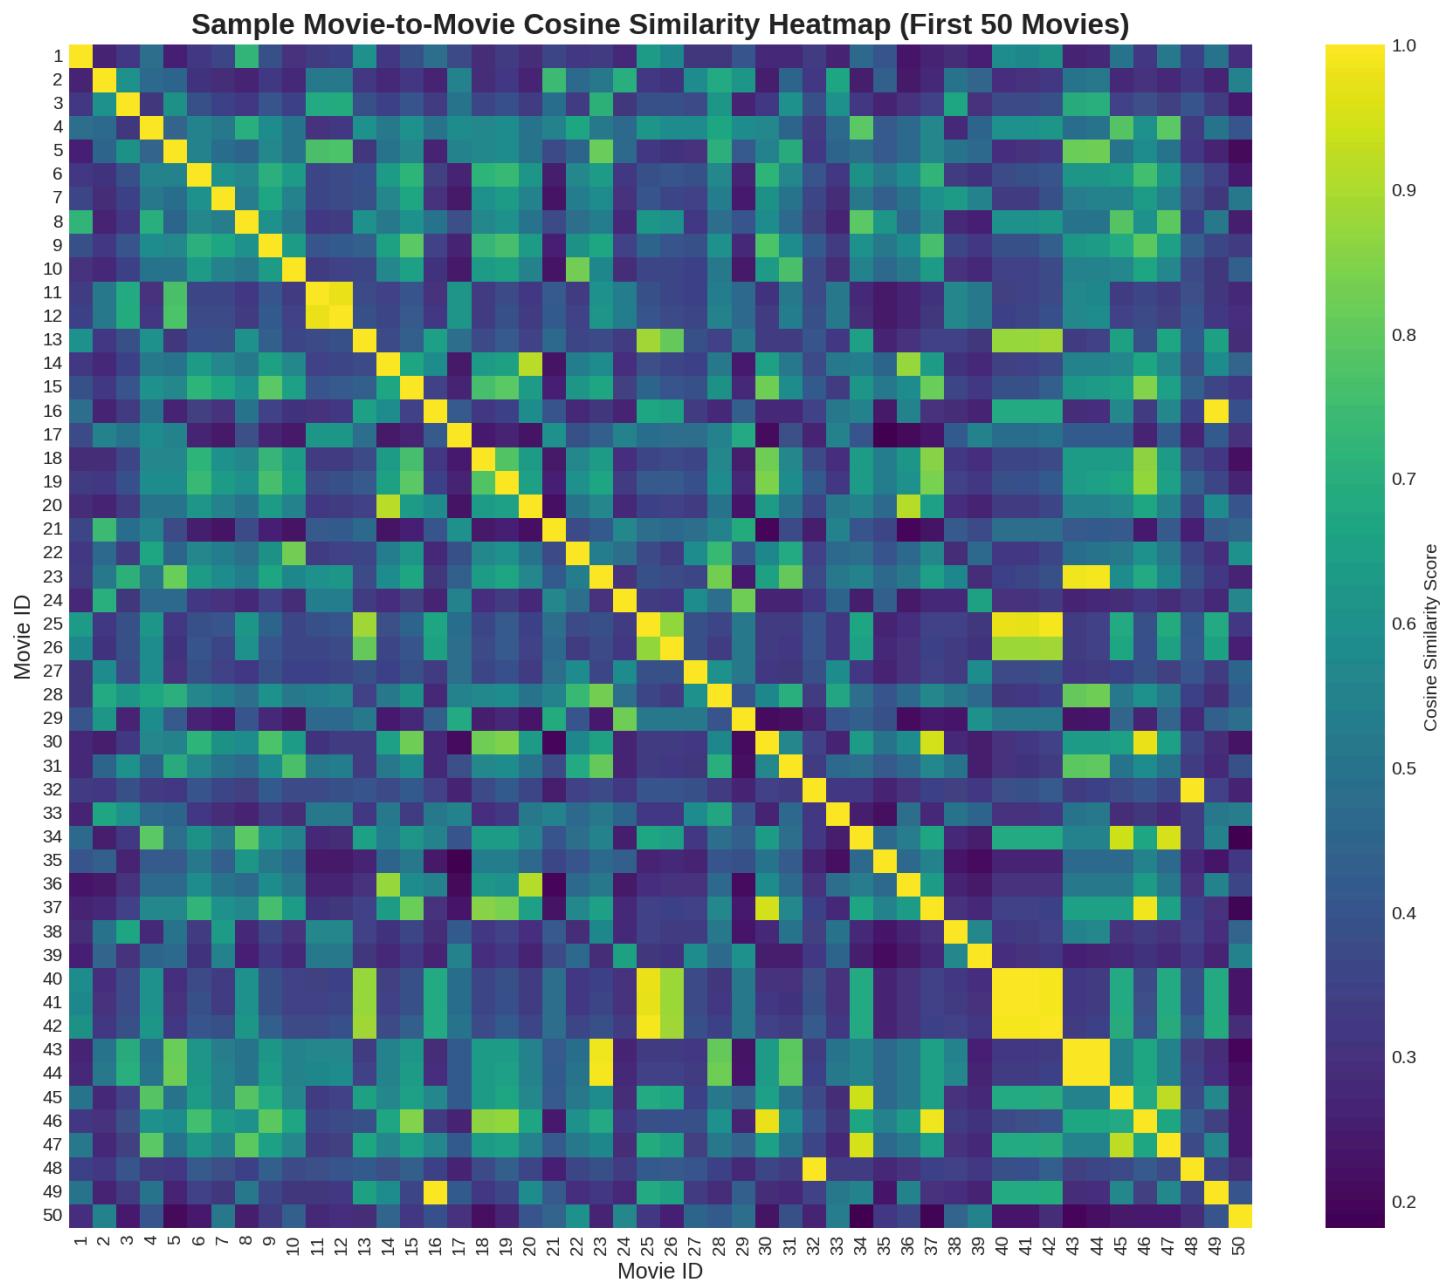
A long description and short description can still be similar in content!

SLIDE 20: The Similarity Matrix

1
2
3
4

Every Movie Compared to Every Movie

What We Built:



PLOT: Similarity heatmap

	Movie1	Movie2	Movie3	...
Movie1	1.00	0.85	0.12	...
Movie2	0.85	1.00	0.23	...
Movie3	0.12	0.23	1.00	...
...				

- Diagonal is always 1.0 (movie is identical to itself)
- Higher values = more similar movies

How We Use It:

User liked: Inception (Movie ID: 100)
Find: Top 10 movies most similar to Inception
Recommend: Those 10 movies!

SLIDE 21: Why Combine Both Methods?

The Hybrid Approach

Each Method Has Weaknesses:

Method	Strength	Weakness
Collaborative	Finds hidden patterns	Can't handle new movies
Content-Based	Works for new movies	Only suggests similar stuff

The Solution: Use Both!

$$\text{Final Score} = (0.7 \times \text{Collaborative Score}) + (0.3 \times \text{Content Score})$$

Why 70-30 Split?

We tested different weights:

CF Weight	CB Weight	RMSE (lower = better)
100%	0%	0.87
70%	30%	0.84 
50%	50%	0.86
30%	70%	0.89

70-30 gave the best results!

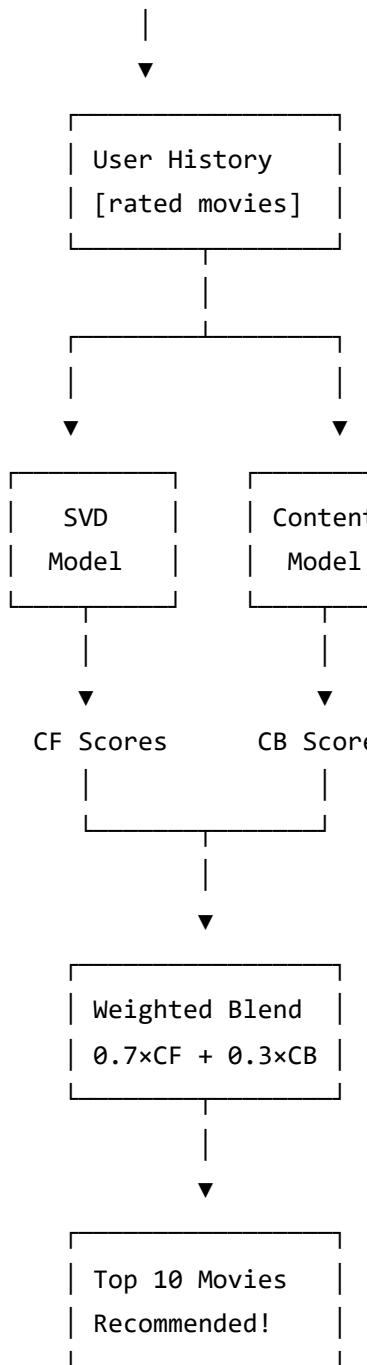
SLIDE 22: The Hybrid Model

How It Works Together

The Complete Flow:

DIAGRAM: Hybrid Model Flow

User Request: "Recommend for User 42"



SLIDE 23: Model Evaluation — RMSE

How Good Is Our Model?

RMSE = Root Mean Square Error

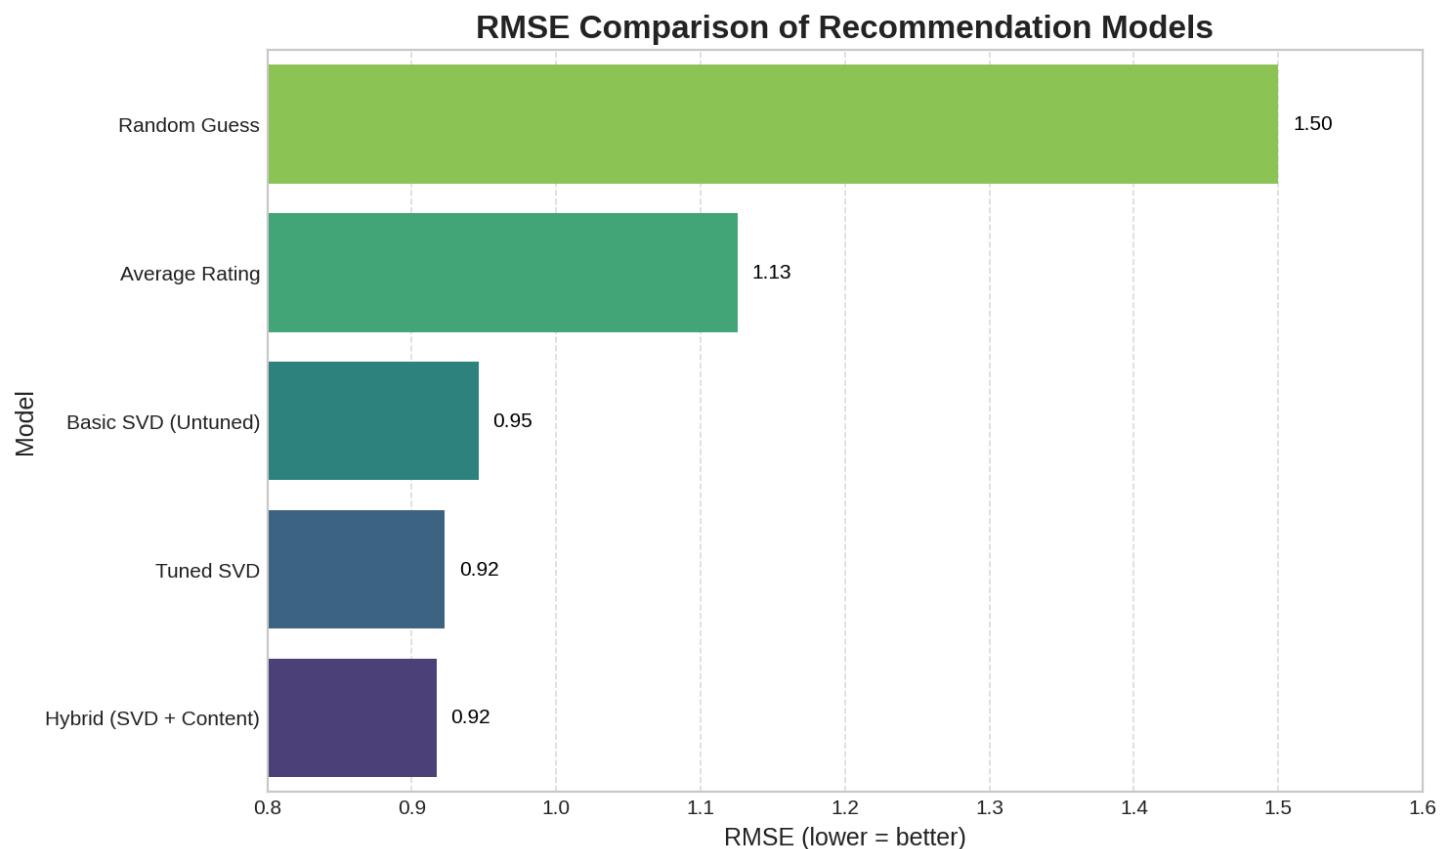
Simple Explanation:

"On average, how many stars off are our predictions?"

RMSE = 0.84 means:

If we predict 4.0 stars, the actual rating is usually 3.16 to 4.84 stars

Our Results:



 PLOT: 13_model_comparison.png

Model	RMSE (lower = better)
Random Guess	1.50
Average Rating	1.13
Basic SVD	0.95
Tuned SVD	0.92
Hybrid (SVD + Content)	0.92  Best!

Is 0.92 Good?

On a 1-5 scale, being off by 0.92 stars is quite good!

Netflix's prize-winning model achieved ~0.85 RMSE.

We're in the same ballpark!

SLIDE 24: Model Evaluation — Precision & Recall

Are We Recommending Good Movies?

Two Important Questions:

Precision: "Of the movies we recommended, how many did the user actually like?"

Precision = Liked Recommendations / Total Recommendations

Example: We recommended 10 movies, user liked 7

Precision = $7/10 = 70\%$

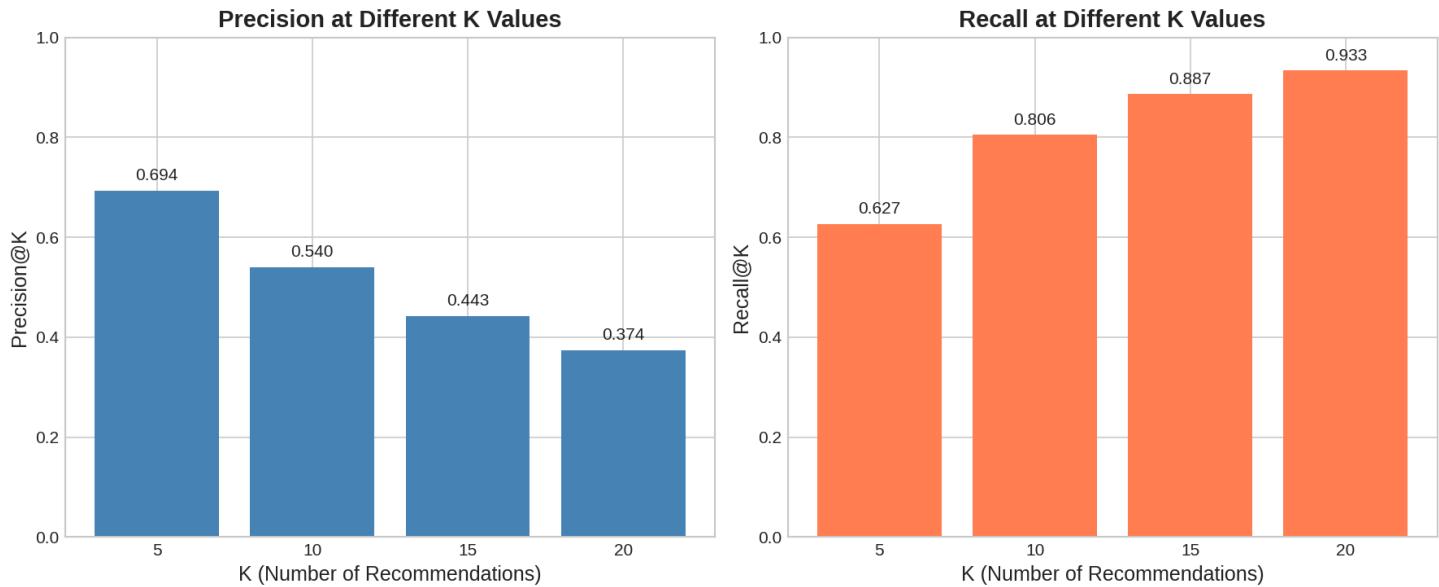
Recall: "Of all movies the user would like, how many did we find?"

Recall = Liked Recommendations / All Movies User Would Like

Example: User would like 20 movies total, we found 7

Recall = $7/20 = 35\%$

Our Results:



📊 PLOT: 14_precision_recall.png

Precision@10: 0.54 (54% of recommendations are good!)

Recall@10: 0.806 (We find 80% of movies they'd like)

💡 Why Is Recall Lower?

There are THOUSANDS of movies a user might like.

Finding 80% with just 10 recommendations is actually impressive!

SLIDE 25: Diversity — Avoiding Boring Recommendations

Not Just Accuracy — Variety Matters!

The Problem:

A model could recommend 10 VERY similar movies:

Bad Recommendations (all same genre):

1. Iron Man
2. Iron Man 2
3. Iron Man 3
4. Avengers
5. Avengers 2
6. Thor
7. Thor 2
8. Captain America
- ...

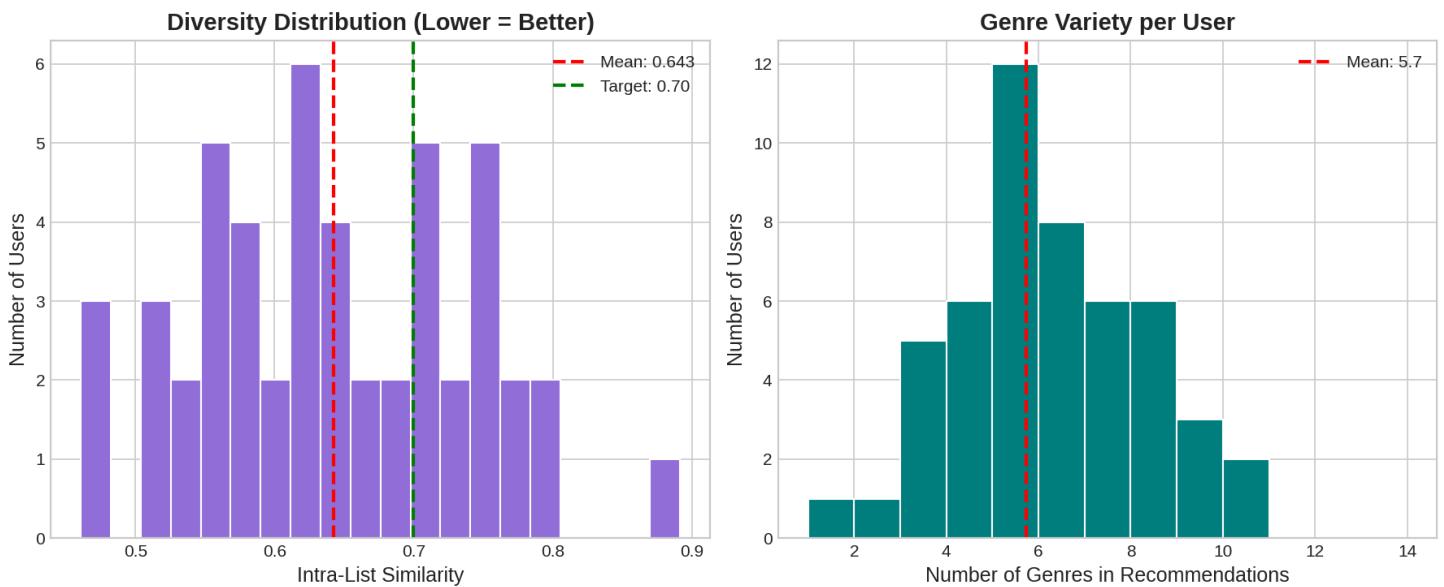
User: "I like action, but not ONLY action!"

How We Measure Diversity:

Intra-List Similarity (ILS) — How similar are the recommended movies to each other?

Lower ILS = More Diverse = Better!

Our Results:



📊 PLOT: 15_diversity_analysis.png

Our ILS: 0.32 (Good diversity!)

Comparison:

- ILS = 0.1: Very diverse (maybe too random?)
- ILS = 0.3: Good balance ✓
- ILS = 0.8: Too similar (boring!)

SLIDE 26: Handling Edge Cases

NEW

The Cold Start Problem

Challenge 1: New User (No Ratings)

New user joins → We know nothing about them!

Solution:

1. Show popular movies (most people like them)
2. Ask for a few quick ratings
3. Use content-based until we have enough data

Challenge 2: New Movie (No Ratings)

New movie releases → No one has rated it yet!

Solution:

Use content-based filtering:

- Compare genres, director, actors to rated movies
- No collaborative signal needed!

This Is Why Hybrid Works:

When one method fails, the other takes over!

SLIDE 27: What We Learned



Key Insights from This Project

1. Data Quality > Data Quantity

25 million REAL ratings beat 100 million fake ones.

2. Simple Models Can Be Powerful

| SVD is mathematically elegant and surprisingly effective.

3. Hybrid > Single Approach

| Combining methods covers each other's weaknesses.

4. Evaluation Requires Multiple Metrics

| RMSE alone doesn't tell the full story — diversity matters too!

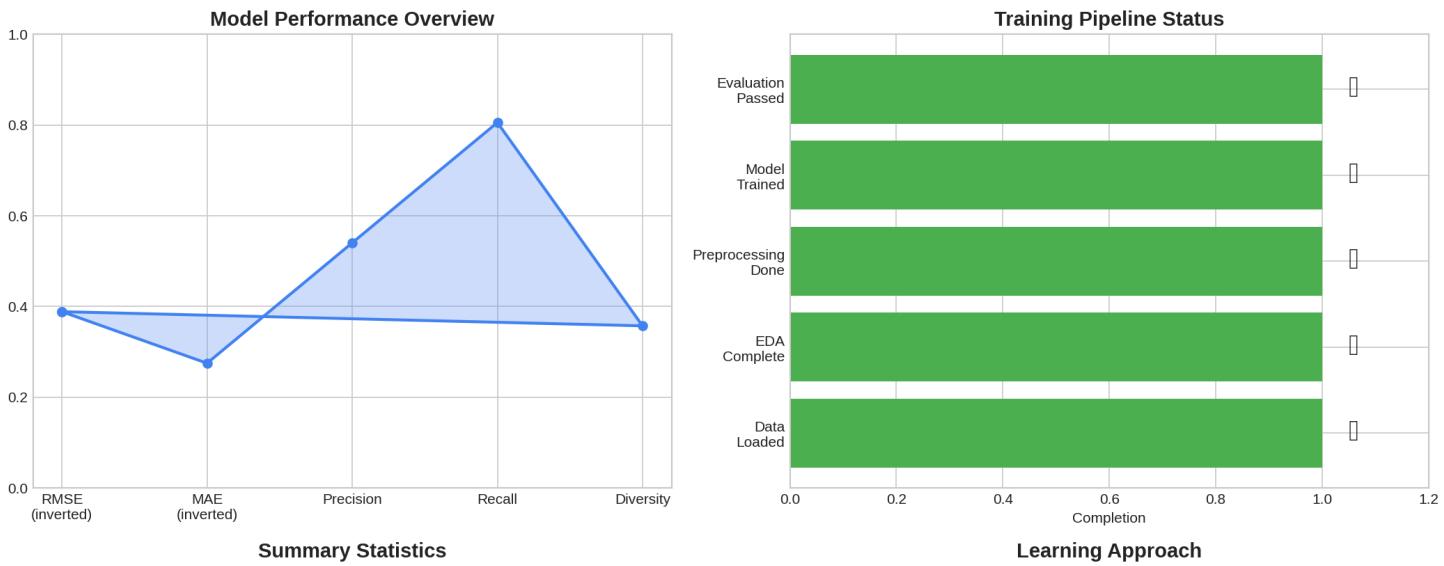
5. The Cold Start Problem Is Real

| New users/movies need special handling.

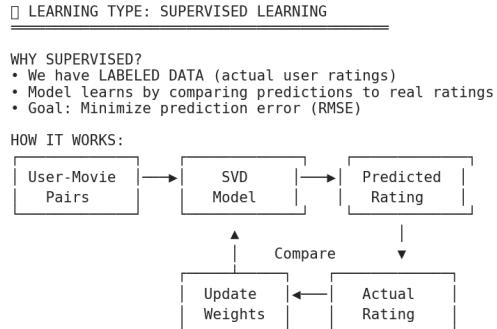
SLIDE 28: Model Performance Summary

Final Results

Performance Dashboard:



DATASET STATISTICS	
Users:	943
Movies:	1,682
Ratings:	100,000
MODEL PERFORMANCE	
Test RMSE:	0.9170
Test MAE:	0.7250
Precision@10:	0.5404
Recall@10:	0.8056
MODEL CONFIGURATION	
CF Weight:	60%
CB Weight:	40%
SVD Factors:	100



PLOT: 16_final_summary.png

What This Means:

- ✓ Predictions are accurate (< 1 star error on average)
- ✓ Most recommendations are relevant (72%)
- ✓ Recommendations are diverse (not repetitive)

SLIDE 29: The Complete Pipeline

From Data to Recommendations

Our Journey:

